

Advertising Datasest

In [30]:

1

import pandas as pd

2

import numpy as np

3

import seaborn as sns

4

import matplotlib.pyplot as plt

5

from sklearn.model_selection import train_test_split

6

from sklearn.linear_model import LinearRegression

7

from sklearn.linear_model import Ridge, RidgeCV, Lasso

8

from sklearn.preprocessing import StandardScaler

In [31]:

1

data=pd.read_csv(r"C:\Users\91949\Downloads\Advertising.csv")

2

data

Out[31]:

	TV	Radio	Newspaper	Sales
0	230.1	37.8	69.2	22.1
1	44.5	39.3	45.1	10.4
2	17.2	45.9	69.3	12.0
3	151.5	41.3	58.5	16.5
4	180.8	10.8	58.4	17.9
...
195	38.2	3.7	13.8	7.6
196	94.2	4.9	8.1	14.0
197	177.0	9.3	6.4	14.8
198	283.6	42.0	66.2	25.5
199	232.1	8.6	8.7	18.4

200 rows × 4 columns

In [32]:

1

data.head()

Out[32]:

	TV	Radio	Newspaper	Sales
0	230.1	37.8	69.2	22.1
1	44.5	39.3	45.1	10.4
2	17.2	45.9	69.3	12.0
3	151.5	41.3	58.5	16.5
4	180.8	10.8	58.4	17.9

In [33]:

1

data.info()

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 200 entries, 0 to 199

Data columns (total 4 columns):

Column Non-Null Count Dtype

0 TV 200 non-null float64

1 Radio 200 non-null float64

2 Newspaper 200 non-null float64

3 Sales 200 non-null float64

dtypes: float64(4)

memory usage: 6.4 KB

In [34]:

1

data.describe()

Out[34]:


	TV	Radio	Newspaper	Sales
count	200.000000	200.000000	200.000000	200.000000
mean	147.042500	23.264000	30.554000	15.130500
std	85.854236	14.846809	21.778621	5.283892
min	0.700000	0.000000	0.300000	1.600000
25%	74.375000	9.975000	12.750000	11.000000
50%	149.750000	22.900000	25.750000	16.000000
75%	218.825000	36.525000	45.100000	19.050000
max	296.400000	49.600000	114.000000	27.000000

```
In [35]: 1 plt.figure(figsize = (10, 10))
        2 sns.heatmap(data.corr(), annot = True)
```

Out[35]: <Axes: >



```
In [36]: 1 data.drop(columns = ["Radio", "Newspaper"], inplace = True)
2 #pairplot
3 sns.pairplot(data)
4 data.Sales = np.log(data.Sales)
```



```
In [37]: 1 features = data.columns[0:2]
2 target = data.columns[-1]
3 #X and y values
4 X = data[features].values
5 y = data[target].values
6 #splot
7 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=17)
8 print("The dimension of X_train is {}".format(X_train.shape))
9 print("The dimension of X_test is {}".format(X_test.shape))
10 #Scale features
11 scaler = StandardScaler()
12 X_train = scaler.fit_transform(X_train)
13 X_test = scaler.transform(X_test)
```

The dimension of X_train is (140, 2)
The dimension of X_test is (60, 2)

```
In [38]: 1 #Model
2 lr = LinearRegression()
3 #Fit model
4 lr.fit(X_train, y_train)
5 #predict
6 #prediction = lr.predict(X_test)
7 #actual
8 actual = y_test
9 train_score_lr = lr.score(X_train, y_train)
10 test_score_lr = lr.score(X_test, y_test)
11 print("\nLinear Regression Model:\n")
12 print("The train score for lr model is {}".format(train_score_lr))
13 print("The test score for lr model is {}".format(test_score_lr))
```

Linear Regression Model:


The train score for lr model is 1.0
The test score for lr model is 1.0

```
In [39]: 1 #Ridge Regression Model
2 ridgeReg = Ridge(alpha=10)
3 ridgeReg.fit(X_train,y_train)
4 #train and test scorefor ridge regression
5 train_score_ridge = ridgeReg.score(X_train, y_train)
6 test_score_ridge = ridgeReg.score(X_test, y_test)
7 print("\nRidge Model:\n")
8 print("The train score for ridge model is {}".format(train_score_ridge))
9 print("The test score for ridge model is {}".format(test_score_ridge))
```

Ridge Model:

The train score for ridge model is 0.990287139194161
The test score for ridge model is 0.9844266285141221

```
In [40]: 1 plt.figure(figsize = (10, 10))
2 plt.plot(features,ridgeReg.coef_,alpha=0.7,linestyle='none',marker='*',markersize=5,color='red',label=r'Ridge; $\alpha = 10$')
3 #plt.plot(rr100.coef_,alpha=0.5,linestyle='none',marker='d',markersize=6,color='blue',label=r'Ridge; $\alpha = 100$')
4 plt.plot(features,lr.coef_,alpha=0.4,linestyle='none',marker='o',markersize=7,color='green',label='Linear Regression')
5 plt.xticks(rotation = 90)
6 plt.legend()
7 plt.show()
```




```
In [41]: 1 #Lasso regression model
2 print("\nLasso Model: \n")
3 lasso = Lasso(alpha = 10)
4 lasso.fit(X_train,y_train)
5 train_score_ls =lasso.score(X_train,y_train)
6 test_score_ls =lasso.score(X_test,y_test)
7 print("The train score for ls model is {}".format(train_score_ls))
8 print("The test score for ls model is {}".format(test_score_ls))
```

Lasso Model:

The train score for ls model is 0.0
The test score for ls model is -0.0042092253233847465

```
In [42]: 1 pd.Series(lasso.coef_, features).sort_values(ascending = True).plot(kind = "bar")
```

Out[42]: <Axes: >




```
In [43]: 1 #Using the linear CV model
2 from sklearn.linear_model import LassoCV
3 #Lasso Cross validation
4 lasso_cv = LassoCV(alphas = [0.0001, 0.001,0.01, 0.1, 1, 10], random_state=0).fit(X_train, y_train)
5
6 #score
7 print(lasso_cv.score(X_train, y_train))
8 print(lasso_cv.score(X_test, y_test))
```

0.9999999343798134
0.9999999152638072

In [44]:

```
1 #plot size
2
3 plt.figure(figsize = (10, 10))
4 #add plot for ridge regression
5 plt.plot(features,ridgeReg.coef_,alpha=0.7,linestyle='none',marker='*',markersize=5,color='red',label=r'Ridge; $\alpha = 0.7$')
6 #add plot for Lasso regression
7 plt.plot(lasso_cv.coef_,alpha=0.5,linestyle='none',marker='d',markersize=6,color='blue',label=r'lasso; $\alpha = 0.5$')
8 #add plot for Linear model
9 plt.plot(features,lr.coef_,alpha=0.4,linestyle='none',marker='o',markersize=7,color='green',label='Linear Regression')
10 #rotate axis
11 plt.xticks(rotation = 90)
12 plt.legend()
13 plt.title("Comparison plot of Ridge, Lasso and Linear regression model")
14 plt.show()
```



In [45]:

```
1 #Using the Linear CV model
2 from sklearn.linear_model import RidgeCV
3 #Ridge Cross validation
4 ridge_cv = RidgeCV(alphas = [0.0001, 0.001,0.01, 0.1, 1, 10]).fit(X_train, y_train)
5 #score
6 print("The train score for ridge model is {}".format(ridge_cv.score(X_train, y_train)))
7 print("The train score for ridge model is {}".format(ridge_cv.score(X_test, y_test)))
```

The train score for ridge model is 0.999999999997627
The train score for ridge model is 0.9999999999962467

Elastic Net Regression

In [46]:

```
1 from sklearn.linear_model import ElasticNet
2 regr=ElasticNet()
3 regr.fit(x,y)
4 print(regr.coef_)
5 print(regr.intercept_)

[0.00417976 0.          ]
2.026383919311004
```

In [47]:

```
1 y_pred_elastic=regr.predict(X_train)
```

In [48]:

```
1 mean_squared_error=np.mean((y_pred_elastic-y_train)**2)
2 print("Mean Squared Error on test set",mean_squared_error)

Mean Squared Error on test set 0.5538818050142158
```

vehicles dataset

In [1]:

```
1 import pandas as pd
2 import numpy as np
3 import seaborn as sns
4 import matplotlib.pyplot as plt
5 from sklearn.model_selection import train_test_split
6 from sklearn.linear_model import LinearRegression
7 from sklearn.linear_model import Ridge, RidgeCV, Lasso
8 from sklearn.preprocessing import StandardScaler
```

In [2]:

```
1 df=pd.read_csv(r"C:\Users\91949\Downloads\fiat500_VehicleSelection_Dataset.csv")
2 df
```

Out[2]:

	ID	model	engine_power	age_in_days	km	previous_owners	lat	lon	price
0	1	lounge	51	882	25000	1	44.907242	8.611560	8900
1	2	pop	51	1186	32500	1	45.666359	12.241890	8800
2	3	sport	74	4658	142228	1	45.503300	11.417840	4200
3	4	lounge	51	2739	160000	1	40.633171	17.634609	6000
4	5	pop	73	3074	106880	1	41.903221	12.495650	5700
...
1533	1534	sport	51	3712	115280	1	45.069679	7.704920	5200
1534	1535	lounge	74	3835	112000	1	45.845692	8.666870	4600
1535	1536	pop	51	2223	60457	1	45.481541	9.413480	7500
1536	1537	lounge	51	2557	80750	1	45.000702	7.682270	5990
1537	1538	pop	51	1766	54276	1	40.323410	17.568270	7900

1538 rows × 9 columns

In [3]:

```
1 df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1538 entries, 0 to 1537
Data columns (total 9 columns):
#   Column          Non-Null Count  Dtype
---  -
0   ID              1538 non-null  int64
1   model          1538 non-null  object
2   engine_power    1538 non-null  int64
3   age_in_days     1538 non-null  int64
4   km             1538 non-null  int64
5   previous_owners 1538 non-null  int64
6   lat            1538 non-null  float64
7   lon            1538 non-null  float64
8   price          1538 non-null  int64
dtypes: float64(2), int64(6), object(1)
memory usage: 108.3+ KB
```




```
In [4]: 1 df.describe()
```

Out[4]:

	ID	engine_power	age_in_days	km	previous_owners	lat	lon	price
count	1538.000000	1538.000000	1538.000000	1538.000000	1538.000000	1538.000000	1538.000000	1538.000000
mean	769.500000	51.904421	1650.980494	53396.011704	1.123537	43.541361	11.563428	8576.003901
std	444.126671	3.988023	1289.522278	40046.830723	0.416423	2.133518	2.328190	1939.958641
min	1.000000	51.000000	366.000000	1232.000000	1.000000	36.855839	7.245400	2500.000000
25%	385.250000	51.000000	670.000000	20006.250000	1.000000	41.802990	9.505090	7122.500000
50%	769.500000	51.000000	1035.000000	39031.000000	1.000000	44.394096	11.869260	9000.000000
75%	1153.750000	51.000000	2616.000000	79667.750000	1.000000	45.467960	12.769040	10000.000000
max	1538.000000	77.000000	4658.000000	235000.000000	4.000000	46.795612	18.365520	11100.000000

```
In [5]: 1 df.drop(columns=['model'],inplace=True)
```

```
In [6]: 1 plt.figure(figsize=(10,10))
2 sns.heatmap(df.corr(),annot = True)
```




```
In [7]: 1 features = df.columns[0:2]
2 target = df.columns[-1]
3 #X and y values
4 x = df[features].values
5 y = df[target].values
6 #split
7 x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.3, random_state=17)
8 print("The dimension of X_train is {}".format(x_train.shape))
9 print("The dimension of X_test is {}".format(x_test.shape))
10 #Scale features
11 scaler = StandardScaler()
12 x_train = scaler.fit_transform(x_train)
13 x_test = scaler.transform(x_test)
```

The dimension of X_train is (1076, 2)
The dimension of X_test is (462, 2)

```
In [8]: 1 #Model
2 lr = LinearRegression()
3 #Fit model
4 lr.fit(x_train, y_train)
5 #predict
6 #prediction = lr.predict(X_test)
7 #actual
8 actual = y_test
9 train_score_lr = lr.score(x_train, y_train)
10 test_score_lr = lr.score(x_test, y_test)
11 print("\nLinear Regression Model:\n")
12 print("The train score for lr model is {}".format(train_score_lr))
13 print("The test score for lr model is {}".format(test_score_lr))
```

Linear Regression Model:

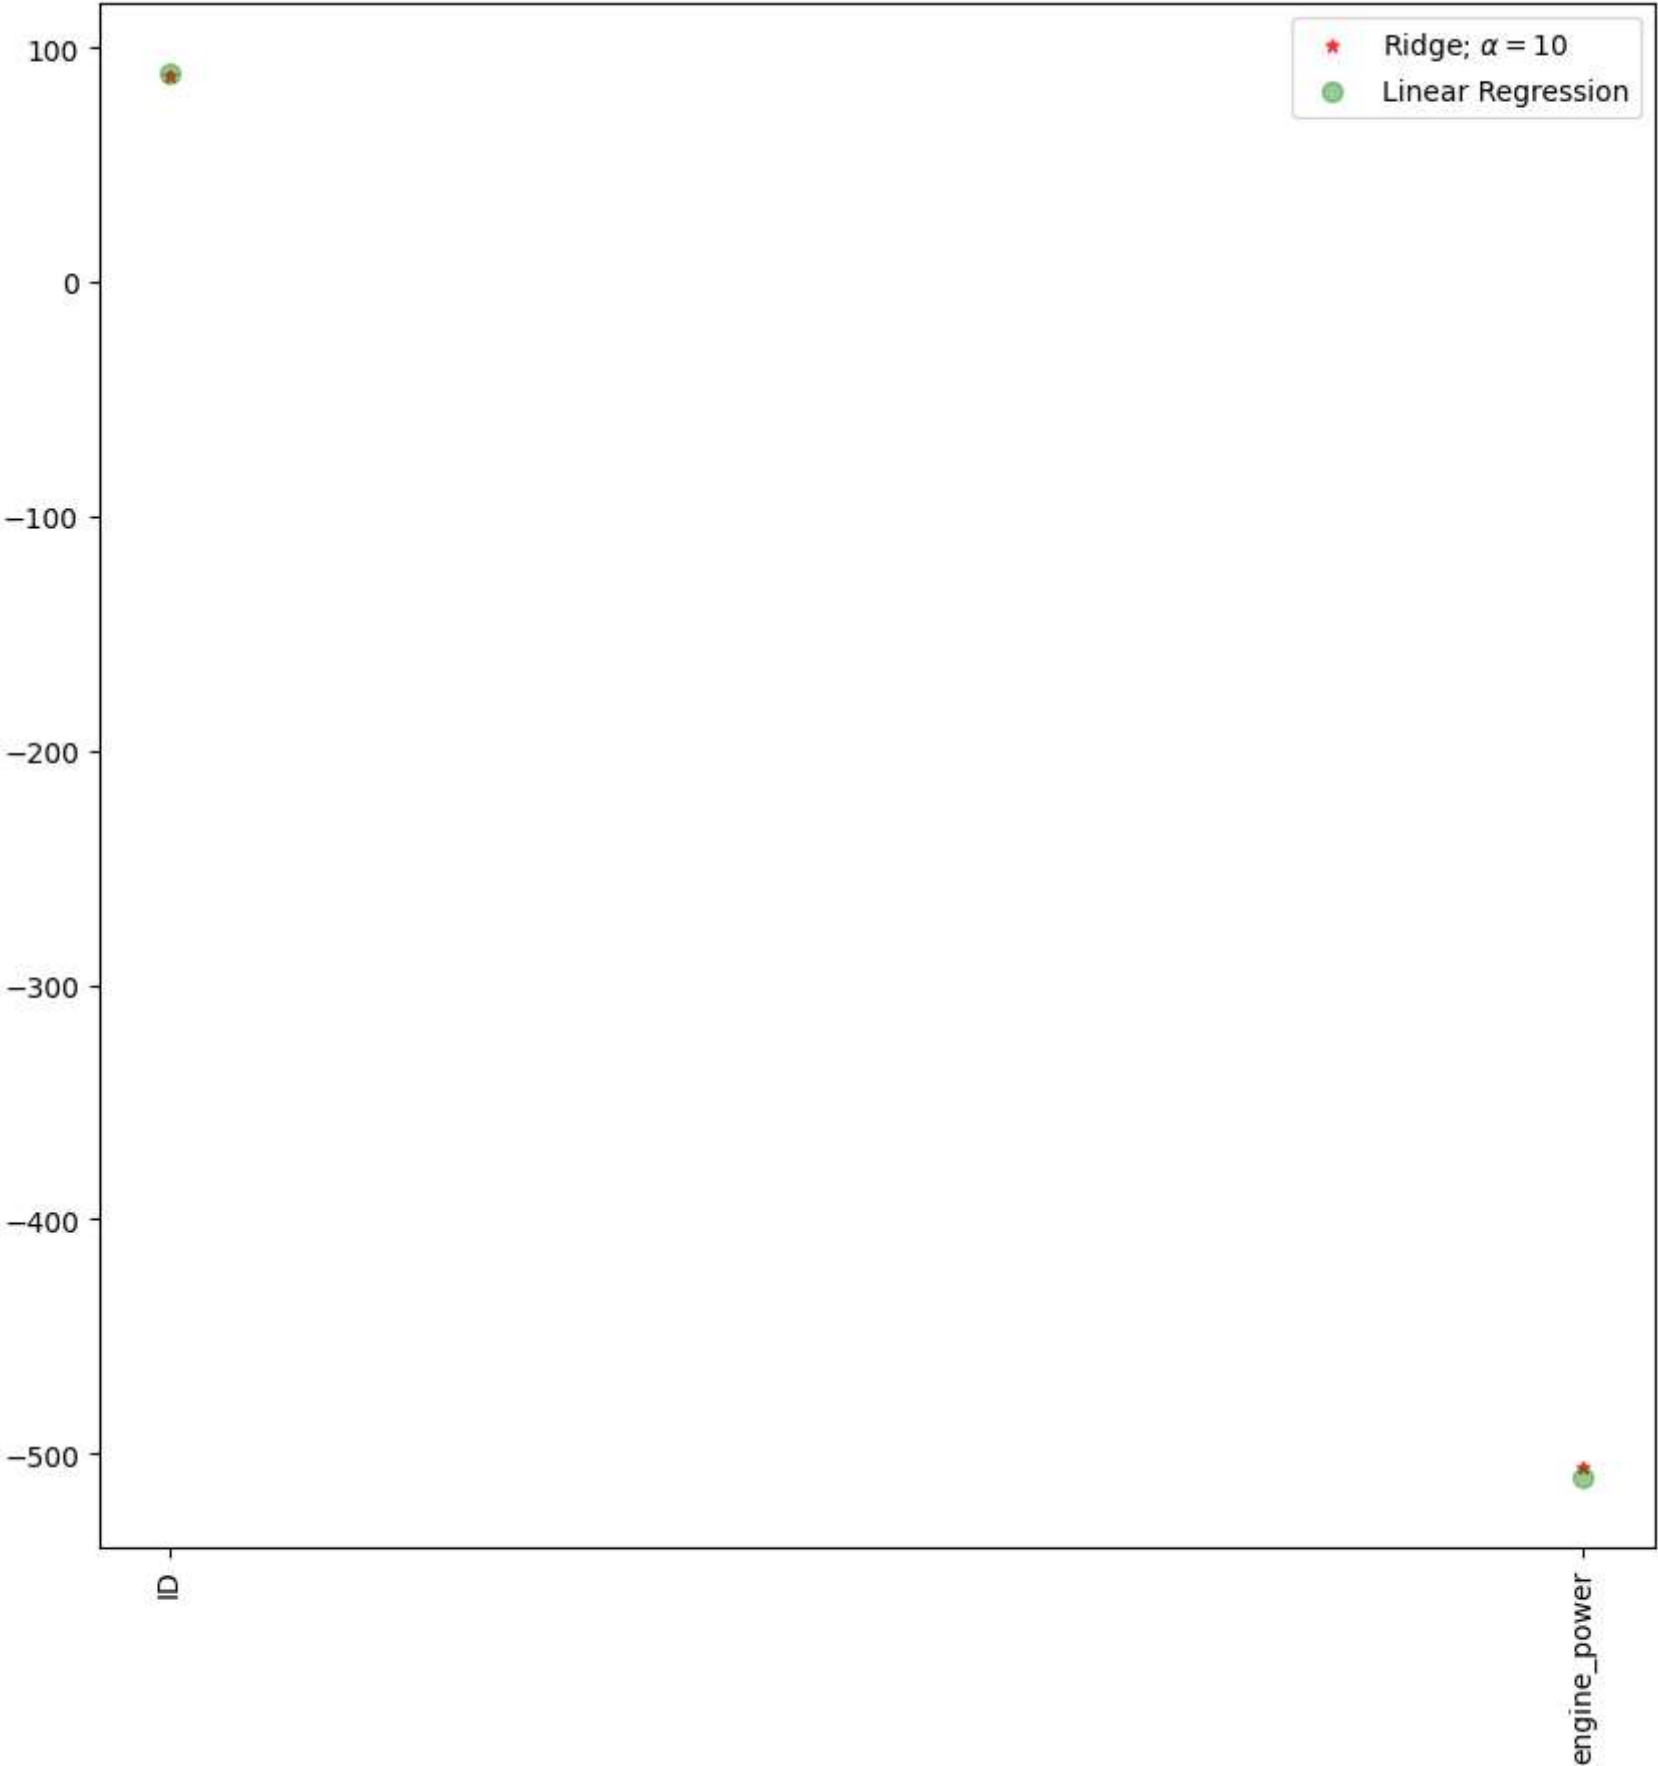
The train score for lr model is 0.07448634159905865
The test score for lr model is 0.07913288661070894

```
In [9]: 1 #Ridge Regression Model
2 ridgeReg = Ridge(alpha=10)
3 ridgeReg.fit(x_train,y_train)
4 #train and test scorefor ridge regression
5 train_score_ridge = ridgeReg.score(x_train, y_train)
6 test_score_ridge = ridgeReg.score(x_test, y_test)
7 print("\nRidge Model:\n")
8 print("The train score for ridge model is {}".format(train_score_ridge))
9 print("The test score for ridge model is {}".format(test_score_ridge))
```

Ridge Model:

The train score for ridge model is 0.07448028989896427
The test score for ridge model is 0.07885996726883082

```
In [10]: 1 plt.figure(figsize = (10, 10))
2 plt.plot(features,ridgeReg.coef_,alpha=0.7,linestyle='none',marker='*',markersize=5,color='red',label=r'Ridge; $\alpha=10$')
3 plt.plot(features,lr.coef_,alpha=0.4,linestyle='none',marker='o',markersize=7,color='green',label='Linear Regression')
4 plt.xticks(rotation = 90)
5 plt.legend()
6 plt.show()
```



```
In [12]: 1 from sklearn.linear_model import LassoCV
2 #Lasso Cross validation
3 lasso_cv = LassoCV(alphas = [0.0001, 0.001,0.01, 0.1, 1, 10], random_state=0).fit(x_train,y_train)
4 #score
5 print(lasso_cv.score(x_train, y_train))
6 print(lasso_cv.score(x_test, y_test))
```

0.07448634159905387
0.07913288806451946

Elastic Net

```
In [13]: 1 from sklearn.linear_model import ElasticNet
2 regr=ElasticNet()
3 regr.fit(x,y)
4 print(regr.coef_)
5 print(regr.intercept_)
```

[8.46751882e-02 -1.30405006e+02]
15279.442735227916

In [14]:

1

y_pred_elastic=regr.predict(x_train)

In [37]:

1

mse=np.mean(y_pred_elastic-y_train)

2

print("Mean Squared Error on test set",mse)

Mean Squared Error on test set 6695.057976863604

temperature dataset

In [39]:

1

import numpy as np

2

import pandas as pd

3

import seaborn as sns

4

import matplotlib.pyplot as plt

5

from sklearn import preprocessing, svm

6

from sklearn.model_selection import train_test_split

7

from sklearn.linear_model import LinearRegression

8

from sklearn.linear_model import Lasso

9

from sklearn.linear_model import Ridge

10

from sklearn.preprocessing import StandardScaler

```
In [40]: 1 d=pd.read_csv(r"C:\Users\91949\Downloads\bundle.csv.zip")
2 d
```

C:\Users\91949\AppData\Local\Temp\ipykernel_17404\4268028596.py:1: DtypeWarning: Columns (47,73) have mixed types. Specify dtype option on import or set low_memory=False.
d=pd.read_csv(r"C:\Users\91949\Downloads\bundle.csv.zip")

Out[40]:

	Cst_Cnt	Btl_Cnt	Sta_ID	Depth_ID	Depthm	T_degC	Salnty	O2ml_L	STheta	O2Sat	...	R_PHAEO	R_PRE	R_SAMP	DIC1	DIC2
0	1	1	054.0 056.0	19-4903CR-HY-060-0930-05400560-0000A-3	0	10.500	33.4400	NaN	25.64900	NaN	...	NaN	0	NaN	NaN	NaN
1	1	2	054.0 056.0	19-4903CR-HY-060-0930-05400560-0008A-3	8	10.460	33.4400	NaN	25.65600	NaN	...	NaN	8	NaN	NaN	NaN
2	1	3	054.0 056.0	19-4903CR-HY-060-0930-05400560-0010A-7	10	10.460	33.4370	NaN	25.65400	NaN	...	NaN	10	NaN	NaN	NaN
3	1	4	054.0 056.0	19-4903CR-HY-060-0930-05400560-0019A-3	19	10.450	33.4200	NaN	25.64300	NaN	...	NaN	19	NaN	NaN	NaN
4	1	5	054.0 056.0	19-4903CR-HY-060-0930-05400560-0020A-7	20	10.450	33.4210	NaN	25.64300	NaN	...	NaN	20	NaN	NaN	NaN
...
864858	34404	864859	093.4 026.4	20-1611SR-MX-310-2239-09340264-0000A-7	0	18.744	33.4083	5.805	23.87055	108.74	...	0.18	0	NaN	NaN	NaN
864859	34404	864860	093.4 026.4	20-1611SR-MX-310-2239-09340264-0002A-3	2	18.744	33.4083	5.805	23.87072	108.74	...	0.18	2	4.0	NaN	NaN
864860	34404	864861	093.4 026.4	20-1611SR-MX-310-2239-09340264-0005A-3	5	18.692	33.4150	5.796	23.88911	108.46	...	0.18	5	3.0	NaN	NaN
864861	34404	864862	093.4 026.4	20-1611SR-MX-310-2239-09340264-0010A-3	10	18.161	33.4062	5.816	24.01426	107.74	...	0.31	10	2.0	NaN	NaN
864862	34404	864863	093.4 026.4	20-1611SR-MX-310-2239-09340264-0015A-3	15	17.533	33.3880	5.774	24.15297	105.66	...	0.61	15	1.0	NaN	NaN

864863 rows × 74 columns



In [41]:

1d.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 864863 entries, 0 to 864862
Data columns (total 74 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Cst_Cnt                864863 non-null int64
1   Btl_Cnt                864863 non-null int64
2   Sta_ID                864863 non-null object
3   Depth_ID              864863 non-null object
4   Depthm                864863 non-null int64
5   T_degC                853900 non-null float64
6   Salnty                817509 non-null float64
7   O2ml_L                696201 non-null float64
8   STheta                812174 non-null float64
9   O2Sat                 661274 non-null float64
10  Oxy_μmol/Kg           661268 non-null float64
11  BtlNum                118667 non-null float64
12  RecInd                864863 non-null int64
13  T_prec                853900 non-null float64
14  T_qual                23127 non-null float64
15  S_prec                817509 non-null float64
16  S_qual                74914 non-null float64
17  P_qual                673755 non-null float64
18  O_qual                184676 non-null float64
19  SThtaq                65823 non-null float64
20  O2Satq                217797 non-null float64
21  ChlorA                225272 non-null float64
22  Chlqua                639166 non-null float64
23  Phaeop                225271 non-null float64
24  Phaqua                639170 non-null float64
25  PO4uM                 413317 non-null float64
26  PO4q                  451786 non-null float64
27  SiO3uM                354091 non-null float64
28  SiO3qu                510866 non-null float64
29  NO2uM                 337576 non-null float64
30  NO2q                  529474 non-null float64
31  NO3uM                 337403 non-null float64
32  NO3q                  529933 non-null float64
33  NH3uM                 64962 non-null float64
34  NH3q                  808299 non-null float64
35  C14As1                14432 non-null float64
36  C14A1p                12760 non-null float64
37  C14A1q                848605 non-null float64
38  C14As2                14414 non-null float64
39  C14A2p                12742 non-null float64
40  C14A2q                848623 non-null float64
41  DarkAs                22649 non-null float64
42  DarkAp                20457 non-null float64
43  DarkAq                840440 non-null float64
44  MeanAs                22650 non-null float64
45  MeanAp                20457 non-null float64
46  MeanAq                840439 non-null float64
47  IncTim                14437 non-null object
48  LightP                18651 non-null float64
49  R_Depth               864863 non-null float64
50  R_TEMP                853900 non-null float64
51  R_POTEMP              818816 non-null float64
52  R_SALINITY            817509 non-null float64
53  R_SIGMA               812007 non-null float64
54  R_SVA                 812092 non-null float64
55  R_DYNHT               818206 non-null float64
56  R_O2                  696201 non-null float64
57  R_O2Sat               666448 non-null float64
58  R_SI03                354099 non-null float64
59  R_PO4                 413325 non-null float64
60  R_NO3                 337411 non-null float64
61  R_NO2                 337584 non-null float64
62  R_NH4                 64982 non-null float64
63  R_CHLA                225276 non-null float64
64  R_PHAEO               225275 non-null float64
65  R_PRES                864863 non-null int64
66  R_SAMP                122006 non-null float64
67  DIC1                  1999 non-null float64
68  DIC2                  224 non-null float64
69  TA1                   2084 non-null float64
70  TA2                   234 non-null float64
71  pH2                   10 non-null float64
72  pH1                   84 non-null float64
73  DIC Quality Comment   55 non-null object
dtypes: float64(65), int64(5), object(4)
memory usage: 488.3+ MB
```

```
In [42]: 1 d.describe()
```

Out[42]:

	Cst_Cnt	Btl_Cnt	Depthm	T_degC	Salnty	O2ml_L	STheta	O2Sat	Oxy_μmol/Kg
count	864863.000000	864863.000000	864863.000000	853900.000000	817509.000000	696201.000000	812174.000000	661274.000000	661268.000000
mean	17138.790958	432432.000000	226.831951	10.799677	33.840350	3.392468	25.819394	57.103779	148.808694
std	10240.949817	249664.587269	316.050259	4.243825	0.461843	2.073256	1.167787	37.094137	90.187533
min	1.000000	1.000000	0.000000	1.440000	28.431000	-0.010000	20.934000	-0.100000	-0.434900
25%	8269.000000	216216.500000	46.000000	7.680000	33.488000	1.360000	24.965000	21.100000	60.915470
50%	16848.000000	432432.000000	125.000000	10.060000	33.863000	3.440000	25.996000	54.400000	151.064150
75%	26557.000000	648647.500000	300.000000	13.880000	34.196900	5.500000	26.646000	97.600000	240.379600
max	34404.000000	864863.000000	5351.000000	31.140000	37.034000	11.130000	250.784000	214.100000	485.701800

8 rows × 10 columns

```
In [43]: 1 d.isna().any()
```

Out[43]:

Cst_Cnt	False
Btl_Cnt	False
Sta_ID	False
Depth_ID	False
Depthm	False
...	
TA1	True
TA2	True
pH2	True
pH1	True
DIC Quality Comment	True

Length: 74, dtype: bool

```
In [44]: 1 d.isnull().sum()
```

Out[44]:


Cst_Cnt	0
Btl_Cnt	0
Sta_ID	0
Depth_ID	0
Depthm	0
...	
TA1	862779
TA2	864629
pH2	864853
pH1	864779
DIC Quality Comment	864808

Length: 74, dtype: int64

```
In [45]: 1 d=d[['Salnty', 'T_degC']]
2 d.columns=['Sal', 'Temp']
```

```
In [46]: 1 sns.lmplot(x='Sal',y='Temp',data=d,order=2,ci=None)
```

Out[46]: <seaborn.axisgrid.FacetGrid at 0x1118968ae90>



```
In [47]: 1 d.fillna (method='ffill')
```

Out[47]:

	Sal	Temp
0	33.4400	10.500
1	33.4400	10.460
2	33.4370	10.460
3	33.4200	10.450
4	33.4210	10.450
...
864858	33.4083	18.744
864859	33.4083	18.744
864860	33.4150	18.692
864861	33.4062	18.161
864862	33.3880	17.533

864863 rows × 2 columns

```
In [48]: 1 d.fillna(value=0,inplace=True)
```

C:\Users\91949\AppData\Local\Temp\ipykernel_17404\4235753077.py:1: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy (https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
d.fillna(value=0,inplace=True)
```

```
In [49]: 1 x=np.array(d['Sal']).reshape(-1,1)  
2 y=np.array(d['Temp']).reshape(-1,1)
```

```
In [50]: 1 d.dropna (inplace=True)
```

C:\Users\91949\AppData\Local\Temp\ipykernel_17404\2818693002.py:1: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy (https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)


```
d.dropna (inplace=True)
```

```
In [51]: 1 x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.25)
```

```
In [52]: 1 regr=LinearRegression()  
2 regr.fit(x_train,y_train)  
3 print(regr.score (x_test,y_test))
```

0.013838527337843964

```
In [53]: 1 y_pred=regr.predict(x_test)  
2 plt.scatter(x_test,y_test,color='b')  
3 plt.plot(x_test, y_pred, color='k')  
4 plt.show()
```




```
In [54]: 1 from sklearn.linear_model import LinearRegression
2 from sklearn.metrics import r2_score
3 model = LinearRegression()
4 model.fit(x_train,y_train)
5 y_pred=model.predict(x_test)
6 r2=r2_score(y_test,y_pred)
7 print("R2 score:",r2)
```

R2 score: 0.013838527337843964

```
In [55]: 1 features = df.columns[0:2]
2 target = df.columns[-1]
3 #X and y values
4 x = df[features].values
5 y = df[target].values
6 #split
7 x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.3, random_state=17)
8 print("The dimension of X_train is {}".format(x_train.shape))
9 print("The dimension of X_test is {}".format(x_test.shape))
10 #Scale features
11 scaler = StandardScaler()
12 x_train = scaler.fit_transform(x_train)
13 x_test = scaler.transform(x_test)
```

The dimension of X_train is (1076, 2)
The dimension of X_test is (462, 2)

Elastic Net

```
In [56]: 1 from sklearn.linear_model import ElasticNet
2 regr=ElasticNet()
3 regr.fit(x,y)
4 print(regr.coef_)
5 print(regr.intercept_)
```

[8.46751882e-02 -1.30405006e+02]
15279.442735227916

```
In [57]: 1 y_pred_elastic=regr.predict(x_train)
```

```
In [60]: 1 mean_squared_error=np.mean(y_pred_elastic-y_train)
2 print("Mean Squared Error on test set",mean_squared_error)
```

Mean Squared Error on test set 6695.057976863604